

Ian Yu

Pragmatic ML Engineer in the Age of Generative AI | 647 972 4689 | ian.yu@arc.com.co | LinkedIn | GitHub

Highlights

- Machine Learning Engineer with 5 years of experience designing and deploying production-grade ML systems in the eCommerce industry. Deployed 4 agentic projects to production and contributed to 2 open research initiatives on Large Language Models (LLMs).
- Proficient in production-grade NLP, agentic workflows, RAG, OCR, MLOps, Geospatial analysis, analytics data engineering, excel at designing-prototype to implement-at-scale transition
- Experienced in consultancy, communication with various levels of stakeholders, product thinking and strategize with users, specialized in utilizing subject matter knowledge with AI

Relevant Experience

Machine Learning Engineer , RezolveAI (formerly Groupby Inc) – Toronto, ON	Aug 2022 – Present
<ul style="list-style-type: none"> Led, designed, and productionized Enrich AI, a major product with agentic workflow that manages product taxonomy, data strategy, and information enrichment, cutting full-service project timelines from weeks to hours. Built autotuning text clustering system with custom loss and outlier handling for efficient AI reasoning reviews, deployed via Ray and Argo on GKE. Lowered orders of magnitude compute while increasing quality by 70% Revamped AI in a patent-pending automotive fitment solution using RAG, achieving >1000 QPS and maintaining <1s latency at 99th percentile. 	
Independent Consultant , Contract – Toronto, ON	Apr 2024 – Present
<ul style="list-style-type: none"> For a RAG ChatBot client with high-fidelity requirements in legal fields, redesigned agentic workflows with 10x lower latency, 3-5x lower cost, while maintaining same output quality Lead, design, and architect for a legal tech company for agentic workflows, legal document extraction, and chatbot with custom-designed memory 	
Data Scientist , Stratica X – Toronto, ON	Apr 2021 – Present
<ul style="list-style-type: none"> For a utility-telecom client, built an end-to-end market analytics system from data engineering to UI to intersect fiber network objects and business directory data with PostGIS. Increased marketing team productivity by 30% and reduced sales cycle time by 30%. 	
Data Strategy Analyst , Groupby Inc – Toronto, ON	Apr 2021 – Jul 2022
<ul style="list-style-type: none"> Created internal packages for client-specific data enrichment analysis, boosting client satisfaction by 50% and reducing revisions by 40%. Automated NLP/ML processes enhancing data observability and collaboration, improving team productivity by 50%. Introduced weakly-supervised text classification system to withhold null tasks from data annotators, saving \$5,500 per month throughout its lifetime. 	

Open Contribution

The BigScience ROOTS Corpus: A 1.6TB Composite Multilingual Dataset [arxiv]	2021 – 2022
<ul style="list-style-type: none"> Led a team of 10 individual contributors to minimize Personally Identifiable Information leakage. Analyzed and defined data filtering parameters to ensure corpus quality, safeness, and ethics NeurIPS 2022, Datasets and Benchmarks Track 	
SantaCoder: don't reach for the stars! [arxiv]	2022 – 2022
<ul style="list-style-type: none"> Advised on data annotation best practices and framework to facilitate PII detection effort for the Stack corpus 	

Technologies

Machine Learning: PyTorch, Transformers, Huggingface, Scikit-learn, SciPy, SpaCy, Snorkel, Vector Databases, Arrows, Polars, SageMaker, Vertex AI, Ray, Kubeflow, MLflow, DSPy
Geospatial: Geopandas, Shapely, turfpy, folium, PostGIS, googlemaps, geojson
Engineering: Python, JavaScript, TypeScript, SQL, NoSQL, Bash, Linux, Docker, Kubernetes, Argo, CI/CD, REST, GraphQL, OpenAPI, Locust, AWS, GCP